

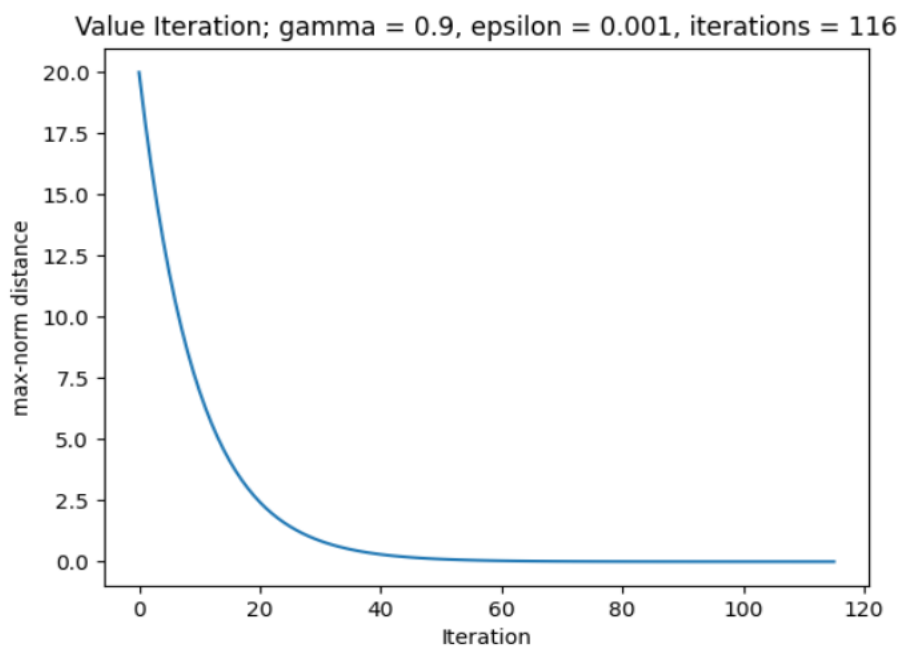
### Part(A): Computing Policies:

1(a,b,c): The Taxi problem has been implemented as a Markov Decision Process(MDP) with states defined by the TaxiPosition, Passenger position, boolean inTaxi(to check if the passenger is inside the taxi). In a 5x5 grid, there are 650 possible states( $25 \times 25$  for not inTaxi + 25 for inTaxi) as when inside, both positions are equal.

The initialization of the problem requires a starting Taxi depot(or position), targetDepot and starting Passenger Depot. Each (state, action) results in a transition to the next state and a reward.

At each location, the taxi can attempt to PICKUP, PUTDOWN deterministically, or go N,E,W,S. Actual motion action is one of (N,E,W,S) with probability 0.85 of being equal to attempted motion action, or 0.05 for other 3 directions. An actual motion action toward a wall leads to no change in state. Default initial State chosen is {TaxiPos: (0, 4), PasPos: (0, 0), inTaxi: False}

2.(a) putting  $\epsilon = 1e-3$  and  $\gamma$ (discount factor) = 0.9, our algorithm converges in 116 iterations.



(b) Epsilon is kept constant =  $1e-3$ , iterations taken to converge are given below:

For gamma = 0.01, Iterations = 3

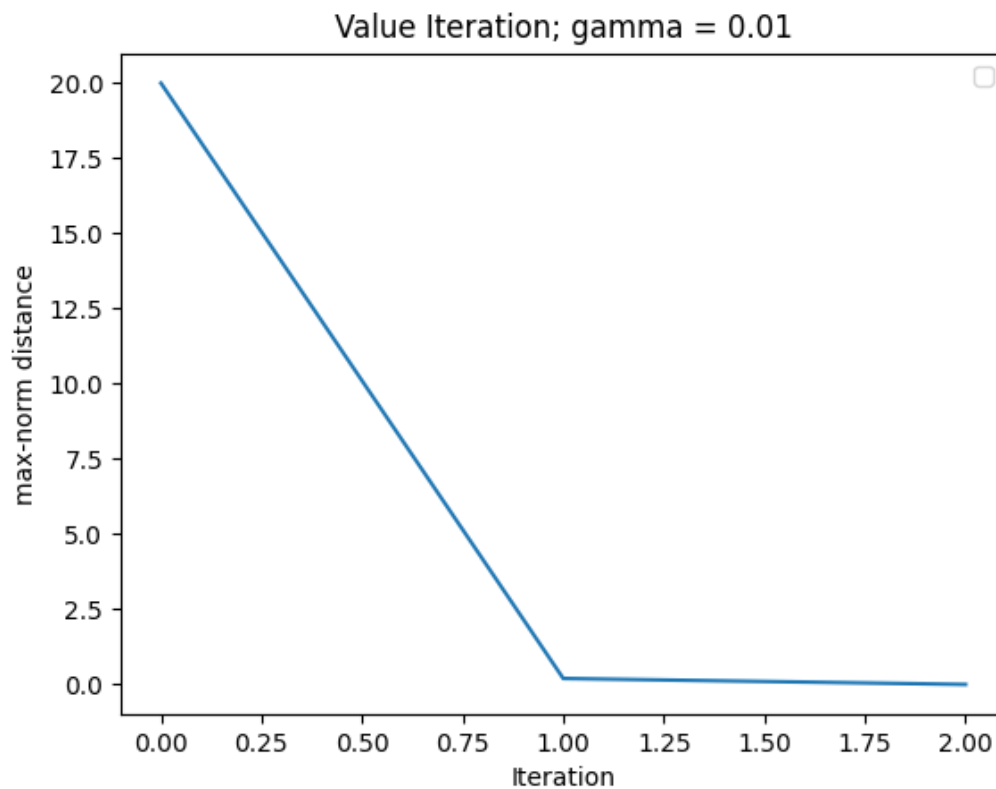
For gamma = 0.1, Iterations = 5



For gamma = 0.5, Iterations = 16



For gamma = 0.8, Iterations = 52

For gamma = 0.99, Iterations = 1444

Above observations can be explained as convergence criterion is  $\delta < \epsilon \cdot (1 - \gamma)$ , where  $\delta$  is the maximum change in utility in consecutive iterations for all states. So, higher gamma would converge at low values of  $\delta$ , and therefore have higher number of iterations. We also see  $\max\_norm\_distance$  decreases as we reach optimal utility values.







(c)

Grid diagram for reference:



We run value iteration with values of  $\gamma = 0.1$ , and  $\gamma = 0.99$ . TargetDepot was fixed as 'G'. All possible combinations of remaining depots were chosen as Initial Passenger and Initial Taxi Position. For initial TaxiDepot = 'R' and initial PassengerDepot = 'Y',  $\gamma = 0.1$ , we get the following sequence:

TaxiStartDepot: ('R', (0, 4)); PassengerStartDepot: ('Y', (0, 0)); DestinationDepot: ('G', (4, 4));  $\gamma = 0.1$

At  $t = 0$ , state  $\rightarrow$  TaxiPos: (0, 4), PasPos: (0, 0), InTaxi: False

Action Taken: S

Reward: -1

At  $t = 1$ , state  $\rightarrow$  TaxiPos: (0, 3), PasPos: (0, 0), InTaxi: False

Action Taken: N

Reward: -1

At  $t = 2$ , state  $\rightarrow$  TaxiPos: (0, 4), PasPos: (0, 0), InTaxi: False

Action Taken: S

Reward: -1

At  $t = 3$ , state  $\rightarrow$  TaxiPos: (0, 3), PasPos: (0, 0), InTaxi: False

Action Taken: N

Reward: -1

At  $t = 4$ , state  $\rightarrow$  TaxiPos: (0, 4), PasPos: (0, 0), InTaxi: False

Action Taken: S

Reward: -1

At  $t = 5$ , state  $\rightarrow$  TaxiPos: (0, 4), PasPos: (0, 0), InTaxi: False

Action Taken: S

Reward: -1

At  $t = 6$ , state  $\rightarrow$  TaxiPos: (0, 3), PasPos: (0, 0), InTaxi: False

Action Taken: N

Reward: -1

At  $t = 7$ , state  $\rightarrow$  TaxiPos: (0, 4), PasPos: (0, 0), InTaxi: False

Action Taken: S

Reward: -1

At  $t = 8$ , state  $\rightarrow$  TaxiPos: (0, 3), PasPos: (0, 0), InTaxi: False

Action Taken: N

Reward: -1

At  $t = 9$ , state  $\rightarrow$  TaxiPos: (1, 3), PasPos: (0, 0), InTaxi: False

Action Taken: N  
Reward: -1

At t = 10, state --> TaxiPos: (1, 4), PasPos: (0, 0), InTaxi: False  
Action Taken: S  
Reward: -1

At t = 11, state --> TaxiPos: (1, 3), PasPos: (0, 0), InTaxi: False  
Action Taken: N  
Reward: -1

At t = 12, state --> TaxiPos: (1, 4), PasPos: (0, 0), InTaxi: False  
Action Taken: S  
Reward: -1

At t = 13, state --> TaxiPos: (1, 3), PasPos: (0, 0), InTaxi: False  
Action Taken: N  
Reward: -1

At t = 14, state --> TaxiPos: (1, 4), PasPos: (0, 0), InTaxi: False  
Action Taken: S  
Reward: -1

At t = 15, state --> TaxiPos: (1, 3), PasPos: (0, 0), InTaxi: False  
Action Taken: N  
Reward: -1

At t = 16, state --> TaxiPos: (1, 4), PasPos: (0, 0), InTaxi: False  
Action Taken: S  
Reward: -1

At t = 17, state --> TaxiPos: (1, 3), PasPos: (0, 0), InTaxi: False  
Action Taken: N  
Reward: -1

At t = 18, state --> TaxiPos: (1, 4), PasPos: (0, 0), InTaxi: False  
Action Taken: S  
Reward: -1

At t = 19, state --> TaxiPos: (1, 3), PasPos: (0, 0), InTaxi: False  
Action Taken: N  
Reward: -1

At t = 20, state --> TaxiPos: (1, 4), PasPos: (0, 0), InTaxi: False  
Action Taken: S  
Reward: -1

Clearly, value iteration algorithm has not converged to an optimal utilities for  $\gamma = 0.1$ , the taxi can be seen oscillating and not moving towards the passenger.

### For $\gamma = 0.99$

TaxiStartDepot: ('R', (0, 4)); PassengerStartDepot: ('Y', (0, 0)); DestinationDepot: ('G', (4, 4));  $\gamma = 0.99$

At t = 0, state --> TaxiPos: (0, 4), PasPos: (0, 0), InTaxi: False  
Action Taken: S  
Reward: -1

At t = 1, state --> TaxiPos: (0, 3), PasPos: (0, 0), InTaxi: False  
Action Taken: S  
Reward: -1

At t = 2, state --> TaxiPos: (0, 2), PasPos: (0, 0), InTaxi: False  
Action Taken: S  
Reward: -1

At t = 3, state --> TaxiPos: (0, 1), PasPos: (0, 0), InTaxi: False  
Action Taken: S

Reward: -1

At t = 4, state --> TaxiPos: (0, 0), PasPos: (0, 0), InTaxi: False  
Action Taken: PICKUP  
Reward: -1

At t = 5, state --> TaxiPos: (0, 0), PasPos: (0, 0), InTaxi: True  
Action Taken: N  
Reward: -1

At t = 6, state --> TaxiPos: (0, 1), PasPos: (0, 1), InTaxi: True  
Action Taken: N  
Reward: -1

At t = 7, state --> TaxiPos: (0, 2), PasPos: (0, 2), InTaxi: True  
Action Taken: E  
Reward: -1

At t = 8, state --> TaxiPos: (1, 2), PasPos: (1, 2), InTaxi: True  
Action Taken: E  
Reward: -1

At t = 9, state --> TaxiPos: (2, 2), PasPos: (2, 2), InTaxi: True  
Action Taken: N  
Reward: -1

At t = 10, state --> TaxiPos: (2, 3), PasPos: (2, 3), InTaxi: True  
Action Taken: E  
Reward: -1

At t = 11, state --> TaxiPos: (3, 3), PasPos: (3, 3), InTaxi: True  
Action Taken: N  
Reward: -1

At t = 12, state --> TaxiPos: (3, 4), PasPos: (3, 4), InTaxi: True  
Action Taken: E  
Reward: -1

At t = 13, state --> TaxiPos: (4, 4), PasPos: (4, 4), InTaxi: True  
Action Taken: PUTDOWN  
Reward: 20

At t = 14, state --> TaxiPos: (4, 4), PasPos: (4, 4), InTaxi: False  
END

We can see at the end of 14 steps, the values have led the taxi to the final state in which the passenger is put down. In this case, coincidentally, random variations are not observed, but another initial state given below shows random noise as well.

TaxiStartDepot: ('R', (0, 4)); PassengerStartDepot: ('B', (3, 0)); DestinationDepot: ('G', (4, 4));  
 $\gamma = 0.99$

At t = 0, state --> TaxiPos: (0, 4), PasPos: (3, 0), InTaxi: False  
Action Taken: S  
Reward: -1

At t = 1, state --> TaxiPos: (0, 3), PasPos: (3, 0), InTaxi: False  
Action Taken: E  
Reward: -1

At t = 2, state --> TaxiPos: (1, 3), PasPos: (3, 0), InTaxi: False  
Action Taken: S  
Reward: -1

At t = 3, state --> TaxiPos: (1, 2), PasPos: (3, 0), InTaxi: False

Action Taken: E  
Reward: -1

At t = 4, state --> TaxiPos: (2, 2), PasPos: (3, 0), InTaxi: False  
Action Taken: E  
Reward: -1

At t = 5, state --> TaxiPos: (3, 2), PasPos: (3, 0), InTaxi: False  
Action Taken: S  
Reward: -1

At t = 6, state --> TaxiPos: (3, 1), PasPos: (3, 0), InTaxi: False  
Action Taken: S  
Reward: -1

At t = 7, state --> TaxiPos: (3, 2), PasPos: (3, 0), InTaxi: False  
Action Taken: S  
Reward: -1

At t = 8, state --> TaxiPos: (3, 1), PasPos: (3, 0), InTaxi: False  
Action Taken: S  
Reward: -1

At t = 9, state --> TaxiPos: (3, 0), PasPos: (3, 0), InTaxi: False  
Action Taken: PICKUP  
Reward: -1

At t = 10, state --> TaxiPos: (3, 0), PasPos: (3, 0), InTaxi: True  
Action Taken: N  
Reward: -1

At t = 11, state --> TaxiPos: (3, 1), PasPos: (3, 1), InTaxi: True  
Action Taken: N  
Reward: -1

At t = 12, state --> TaxiPos: (3, 2), PasPos: (3, 2), InTaxi: True  
Action Taken: N  
Reward: -1

At t = 13, state --> TaxiPos: (3, 3), PasPos: (3, 3), InTaxi: True  
Action Taken: N  
Reward: -1

At t = 14, state --> TaxiPos: (3, 4), PasPos: (3, 4), InTaxi: True  
Action Taken: E  
Reward: -1

At t = 15, state --> TaxiPos: (4, 4), PasPos: (4, 4), InTaxi: True  
Action Taken: PUTDOWN  
Reward: 20

At t = 16, state --> TaxiPos: (4, 4), PasPos: (4, 4), InTaxi: False  
END


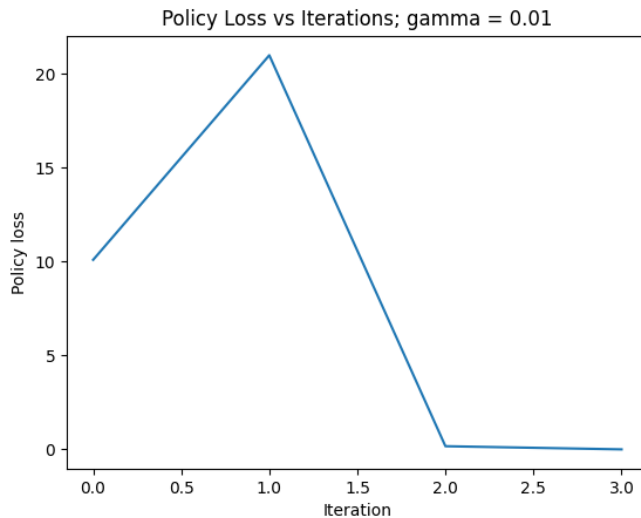
From t=6 to t=7, non-intended action occurs showing stochasticity in the environment model. Again, we see the taxi successfully picks up and drops passenger despite random variations being experienced.

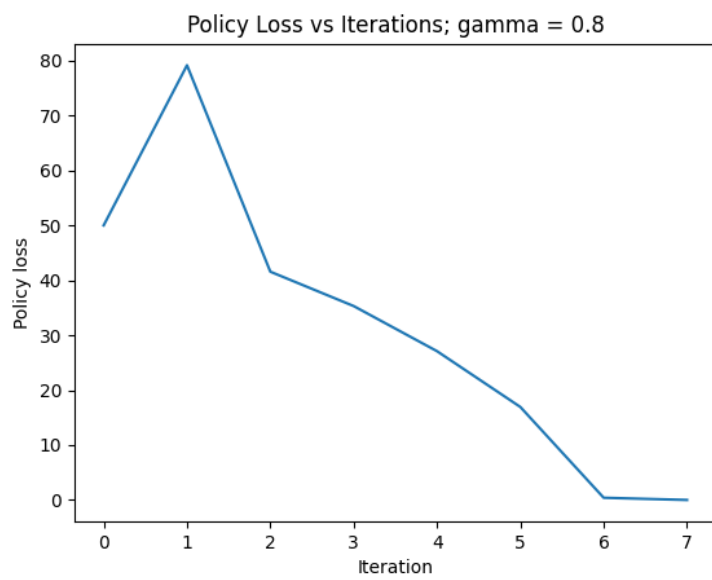
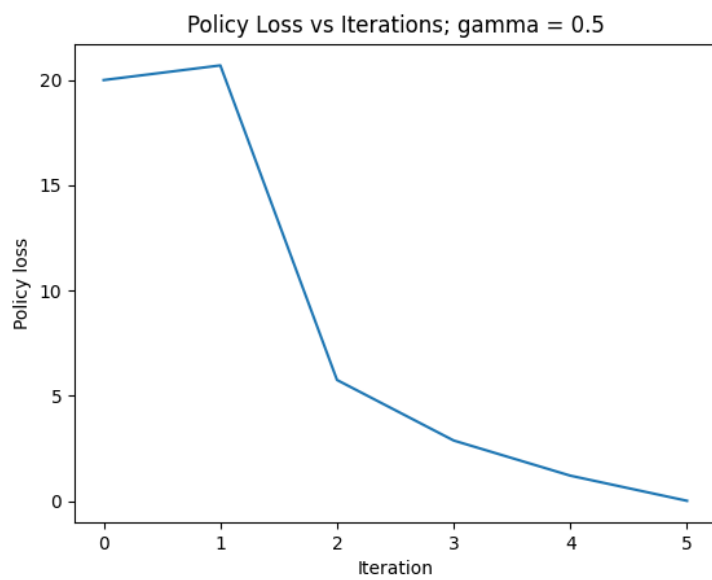
3. (a) For policy iteration, we apply linear algebra/iterative policy evaluation followed by policy improvement. During experimentation, we found linear algebra solvers to be faster for higher values of gamma as calculation of utilities in policy evaluation is done in a single step while it takes multiple iterations during iterative. For higher values of gamma, iterations are more due to convergence being later(as explained in 2).




(b) We plotted policy iteration vs loss for multiple values of  $\gamma = [0.01, 0.1, 0.5, 0.8, 0.9]$ . For all values of  $\gamma$ , policy iteration converged in 5-6 iterations, much lower than value iterations. However, we observe that policy loss values for  $\gamma = 0.99$  were much higher indicating a steep change in randomized policy. This can be explained by the higher contribution of previous values due to large discount factor leading to large change in policy values.

Plots for policy\_loss vs iterations are:








Plots observed are non-steady at different steps, movement of policy is different, hence graph has irregular descent, however policy still finally converges.


## Reinforcement Learning:

2. We set the destination as (4,4). As the following plots show, introducing a decaying exploration rate slows down convergence. We see that Q-Learning with a fixed rate converges the fastest, and Q-Learning with a decaying rate converges the slowest. All of these roughly converge to the same value of average reward over multiple runs, except SARSA with decaying rate, that converges to a negative value. Their final reward values are:


- A. Q-Learning, Fixed : 1.637
- B. Q-Learning, Decaying: 1.04
- C. SARSA, Fixed: 2.06
- D. SARSA, Decaying: -24.29




Reward vs Number of Training Episodes (Q-Learning)  
Epsilon=0.1 fixed, gridSize=5  
alpha=0.25, gamma=0.99




Reward vs Number of Training Episodes (SARSA)  
Epsilon=0.1 decaying, gridSize=5  
alpha=0.25, gamma=0.99




Reward vs Number of Training Episodes (SARSA)  
Epsilon=0.1 fixed, gridSize=5  
alpha=0.25, gamma=0.99




4. Keeping a very low exploration rate means that we are stuck with the initial Q Value estimates, and the resulting policy is nearly random. This results in high stochasticity and poor rewards. On the other hand, keeping a very high exploration rate also means that most of our actions are randomly selected and thus results in high stochasticity of the observed rewards. When we vary the learning rate, we see that a high learning rate places too much importance on the just recently acquired sample over the previous estimates we have learnt. On the other hand, a low learning rate keeps us stuck with our initial estimates too much, and doesn't allow us to use the new samples to update our estimates and learn.




Reward vs Number of Training Episodes (Q-Learning)  
Epsilon=0.05 fixed, gridSize=5  
alpha=0.1, gamma=0.99




Reward vs Number of Training Episodes (Q-Learning)  
Epsilon=0.1 fixed, gridSize=5  
alpha=0.1, gamma=0.99



Reward vs Number of Training Episodes (Q-Learning)  
Epsilon=0.5 fixed, gridSize=5  
alpha=0.1, gamma=0.99




Reward vs Number of Training Episodes (Q-Learning)  
Epsilon=0.9 fixed, gridSize=5  
alpha=0.1, gamma=0.99






Changing alpha:


Reward vs Number of Training Episodes (Q-Learning)  
Epsilon=0.1 fixed, gridSize=5  
alpha=0.1, gamma=0.99




Reward vs Number of Training Episodes (Q-Learning)  
Epsilon=0.1 fixed, gridSize=5  
alpha=0.2, gamma=0.99



Reward vs Number of Training Episodes (Q-Learning)  
Epsilon=0.1 fixed, gridSize=5  
alpha=0.3, gamma=0.99



Reward vs Number of Training Episodes (Q-Learning)  
Epsilon=0.1 fixed, gridSize=5  
alpha=0.4, gamma=0.99



Reward vs Number of Training Episodes (Q-Learning)  
Epsilon=0.1 fixed, gridSize=5  
alpha=0.5, gamma=0.99

